

Sandeep Upreti — Job Search & Personal Branding Toolkit

Target Role: Data Engineer (Azure / Databricks) — Alian Technologies | Remote | ₹12-15 LPA

Contact: sandeepupreti274@gmail.com • +91-9897118186 **LinkedIn:**

linkedin.com/in/sandeep-upreti-83a2a4150 • **GitHub:** github.com/supreti115631 •

Portfolio: sandeepupretiportfolio.netlify.app

Fit Snapshot: You match every *mandatory* Azure pillar of this role — Data Factory, Databricks, Data Lake Storage, ETL/ELT, Spark, Python, SQL, data modeling, and data governance — and you bring Solution Architect-level depth plus 6 Microsoft/Databricks certifications to a Data Engineer requirement. You are a strong-to-overqualified candidate; the messaging below positions you as a hands-on engineer who also thinks like an architect.

SECTION 1: ATS-Friendly Cover Letter (~410 words)

Dear Hiring Manager,

I'm applying for the Data Engineer position at Alian Technologies. Your requirement reads almost exactly like the work I do every day: designing scalable ETL/ELT pipelines on Azure and Databricks, modeling enterprise data, and enforcing governance and quality across cloud data platforms. With 8+ years building cloud data solutions — and six Azure and Databricks certifications, including Databricks Certified Data Engineer Professional — I can contribute from week one.

In my current role as Data Solution Architect and Azure Databricks Engineer at EXL Service, I architected an end-to-end migration from MS SQL Server to Azure Databricks for a large healthcare provider, implementing medallion architecture with Delta Lake and Unity Catalog for enterprise governance. I designed metadata-driven ETL pipelines using Azure Data Factory and Databricks that integrate four distinct source systems — Epic EHR, Microsoft Dynamics 365, Workday, and Five9 — into a single, governed lakehouse. This is the precise stack your role calls for: Data Factory, Databricks, Data Lake Storage, ELT orchestration, and Spark-based processing in Python and PySpark.

Your emphasis on data modeling and governance aligns closely with my experience. I've built enterprise data models across clinical, operational, and financial domains while maintaining HIPAA-aligned data privacy and regulatory compliance — discipline that transfers directly to any environment where data quality and security are non-negotiable. Earlier, as an RCM Subject Matter Expert, I designed custom SQL data models for revenue tracking, AR, and denial management, and optimized SQL Server data warehousing for performance.

Beyond pipeline engineering, I bring strong programming skills in Python, PySpark, Spark SQL, and SQL, plus my Microsoft Certified: Azure Administrator Associate credential, which grounds me in Azure identity, security, and platform administration. I'm also actively expanding into Gen AI and Agentic AI on the Databricks lakehouse, so I bring a forward-looking view of where data platforms are heading.

I'm drawn to this role because it's a focused, build-heavy engineering mandate on the exact Azure ecosystem I've specialized in — and because the remote structure lets me deliver impact without geographic constraint. I'm confident I can design, optimize, and operate the kind of resilient, governed pipelines Alian Technologies needs.

I'd welcome the chance to discuss how my Azure Databricks expertise can support your data platform goals. Thank you for considering my application.

Sincerely, Sandeep Upreti

SECTION 2: Recruiter Outreach Email

Subject: Azure Databricks Data Engineer (8+ yrs, 6 certs) — interested in your Data Engineer opening

Hi [Recruiter Name],

I came across the Data Engineer role at Alian Technologies and wanted to introduce myself directly, as it's a near-exact match to my background.

I'm a Data Solution Architect / Azure Databricks Engineer with 8+ years of experience building ETL/ELT pipelines on Azure Data Factory, Databricks, and Data Lake Storage. Quick highlights:

- Architected a full MS SQL Server → Azure Databricks migration using medallion architecture, Delta Lake, and Unity Catalog
- Built metadata-driven pipelines integrating Epic EHR, D365, Workday, and Five9
- Strong in Python, PySpark, Spark SQL, and SQL, with data modeling and governance built in
- Hold 6 certifications: Databricks Data Engineer Professional + Associate, Databricks GenAI Engineer, Databricks Data Analyst, Databricks Solution Architect Champion, and Microsoft Azure Administrator Associate

I'd love to learn more about the role and share how my experience maps to it. Are you open to a quick call this week?

Resume attached. Thanks for your time.

Best, Sandeep Upreti +91-9897118186 • [linkedin.com/in/sandeep-upreti-83a2a4150](https://www.linkedin.com/in/sandeep-upreti-83a2a4150)

SECTION 3: Hiring Manager Email

Subject: Data Engineer application — Azure Databricks specialist ready to build

Hi [Hiring Manager Name],

I'm reaching out about the Data Engineer role on your team. I build the exact systems your job description outlines — scalable ETL/ELT pipelines on Azure Data Factory and Databricks, governed by Delta Lake and Unity Catalog — and I'd like to make the case briefly.

In my current role at EXL Service, I led the migration of a healthcare provider's analytics platform from MS SQL Server to Azure Databricks, implementing medallion architecture and metadata-driven ingestion across four enterprise source systems. The work demanded everything in your mandatory list: Data Factory orchestration, Spark/PySpark processing, Python and SQL, dimensional data modeling, and strict data governance and quality controls under HIPAA.

Two things I'd bring beyond the spec:

1. **Architect's lens on an engineer's job** — I design pipelines to scale, stay observable, and stay cheap (I work with Azure FinOps for cost control), not just to run once.
2. **Governance discipline** — operating in regulated healthcare means data quality and security aren't afterthoughts for me; they're how I build by default.

I'd welcome 15 minutes to understand your current data challenges and where I could plug in fastest. Would [day/time] work?

Best regards, Sandeep Upreti sandeepupreti274@gmail.com • +91-9897118186

SECTION 4: LinkedIn Connection Request (≤300 characters)

Hi [Name], I'm an Azure Databricks Data Engineer (8+ yrs, 6 Azure/Databricks certs) specializing in ETL/ELT, Delta Lake, and Unity Catalog governance. I noticed Alian Technologies' Data Engineer opening and your work — would love to connect and learn more. Thanks!

(Character count: ~285 — safely under the 300 limit.)

SECTION 5: Referral Request Message

Hi [Name],

Hope you're doing well! I noticed Alian Technologies is hiring a Data Engineer, and it's a strong match for my background — I'm an Azure Databricks Engineer with 8+ years in

ETL/ELT pipeline design, Delta Lake, and data governance, plus the Databricks Data Engineer Professional cert.

Since you're connected to the team, would you be open to referring me or pointing me to the right person? I'd be glad to send my resume and a short summary of how I fit so it's easy for you to pass along. No pressure at all if it's not a fit — I appreciate you either way.

Thanks so much, Sandeep

SECTION 6: Follow-Up Email (send 5 days after no response)

Subject: Re: Data Engineer application — following up

Hi [Name],

I wanted to gently follow up on my application for the Data Engineer role I sent last week. I remain very interested — the Azure Databricks and ETL/ELT focus lines up directly with what I do day to day, including a recent end-to-end SQL Server-to-Databricks migration using medallion architecture and Unity Catalog.

If it's helpful, I'm happy to share a brief walkthrough of a relevant pipeline project or answer any questions about my fit. Is there a good time this week for a quick conversation?

Thanks again for your consideration — I look forward to hearing from you.

Best, Sandeep Upreti +91-9897118186 • [linkedin.com/in/sandeep-upreti-83a2a4150](https://www.linkedin.com/in/sandeep-upreti-83a2a4150)

SECTION 7: 30-Second Professional Introduction

"I'm Sandeep Upreti, a Data Solution Architect and Azure Databricks Engineer with 8+ years building cloud data platforms. I specialize in designing scalable ETL/ELT pipelines on Azure Data Factory and Databricks — most recently leading a full migration from SQL Server to a Databricks lakehouse using medallion architecture, Delta Lake, and Unity Catalog for governance. I work primarily in Python, PySpark, and SQL, and I've spent the last several years in healthcare analytics where data quality and HIPAA compliance are critical. I hold six Azure and Databricks certifications, including the Databricks Data Engineer Professional. What I enjoy most is turning messy, multi-source data into governed, reliable platforms that teams can actually trust to make decisions."

(Reads in ~28–30 seconds at a natural pace.)

SECTION 8: Top 10 Job Titles Best Suited for Your Profile

1. **Data Engineer (Azure / Databricks)** — direct match to this role

- 2. **Senior Data Engineer** — your depth and certs support a senior title
 - 3. **Azure Databricks Engineer**
 - 4. **Data Solution Architect** — your current title; lateral or aspirational targets
 - 5. **Lakehouse / Data Platform Engineer**
 - 6. **ETL / ELT Developer (Azure)**
 - 7. **Cloud Data Engineer**
 - 8. **Analytics Engineer**
 - 9. **Data Platform Modernization Lead / Migration Specialist**
 - 10. **Healthcare Data Engineer / Healthcare Analytics Engineer** (domain-specialist track)
-

SECTION 9: Key Strengths Recruiters Will Notice

- **End-to-end Azure stack ownership** — Data Factory, Databricks, ADLS, and Delta Lake, used together on a real production migration, not in isolation.
 - **Six relevant certifications** — including Databricks Data Engineer Professional and Microsoft Azure Administrator Associate, which independently validate your skills.
 - **Proven platform modernization** — you led an actual SQL Server → Databricks lakehouse migration, which is exactly what many employers are hiring to do right now.
 - **Governance built in** — Unity Catalog, HIPAA-aligned compliance, and data quality discipline make you low-risk for regulated or security-conscious teams.
 - **Strong, modern programming core** — Python, PySpark, Spark SQL, and SQL across architecture and hands-on engineering.
 - **Career progression** — Data Analyst → Senior Analyst/SME → Solution Architect shows growth, ownership, and increasing scope at a single reputable firm (EXL Service).
 - **Multi-source integration experience** — Epic EHR, D365, Workday, and Five9 demonstrate you can handle heterogeneous, enterprise-grade data.
 - **Forward-looking** — active in Gen AI / Agentic AI on the lakehouse, signaling you won't go stale.
-

SECTION 10: Skill Gap Analysis

Current Strengths (directly required by the JD)

JD Requirement	Your Evidence
Azure Data Factory	Built metadata-driven pipelines on ADF

JD Requirement	Your Evidence
Azure Databricks	Core specialization; led platform migration
Azure Data Lake Storage	Part of your stated cloud platform stack
ETL/ELT	Designed/deployed across multiple engagements
Spark	Apache Spark, PySpark, Spark SQL
Python & SQL	Primary programming languages
Data modeling	Enterprise models across clinical/operational/financial domains
Data governance & quality	Unity Catalog, HIPAA compliance, data privacy

Missing or Not-Yet-Demonstrated Skills (named in the JD)

- **Azure Synapse Analytics** — JD lists it as mandatory; your resume centers on Databricks instead. *Highest-priority gap.*
- **Kafka** — listed as mandatory; not on your resume.
- **Hadoop ecosystem** — listed as mandatory; not explicitly shown.
- **Cloud security specifics** — Azure AD, Key Vault, VPCs, encryption (your Azure Administrator cert partially covers this; make it explicit).
- **ETL tools (preferred):** Apache Airflow, Informatica, Talend, Apache NiFi — you use ADF/Databricks Workflows/Control-M instead.
- **Java / Scala / R (preferred)** — Python-focused; no JVM languages shown.
- **On-prem cluster management (preferred):** Cloudera / Hortonworks — not shown.
- **Linux/Windows server & networking (preferred)** — not explicitly evidenced.

Recommended Learning Roadmap (priority order)

1. **Azure Synapse Analytics (1–2 weeks):** It's a mandatory line item. Do Microsoft Learn's Synapse path; build one small Synapse pipeline so you can speak to it. Even better, frame your Databricks experience as "modern lakehouse alternative to Synapse" — you can bridge the gap conversationally while you upskill.
2. **Kafka fundamentals (1 week):** Learn producers/consumers, topics, partitions, and how streaming feeds a Databricks pipeline (Structured Streaming + Kafka source). One hands-on demo is enough to interview well.
3. **Make your cloud-security story explicit (a few days):** You already hold the Azure Administrator cert — write resume/talking points that name Azure AD, Key Vault secret scopes in Databricks, and encryption at rest/in transit.
4. **Hadoop ecosystem literacy (3–4 days):** You don't need deep Hadoop ops; understand HDFS, Hive, and how Spark relates, so you can speak to it credibly.

5. **Apache Airflow (1 week):** The most marketable orchestration tool to add alongside ADF; map your existing orchestration experience onto Airflow DAG concepts.
6. **Scala basics (optional, ongoing):** Since Spark is Scala-native, light Scala fluency is a strong differentiator if you have time after the above.

Strategy note: This role lists Synapse, Kafka, and Hadoop as mandatory but is priced and scoped as a single Data Engineer position — many such postings flex on "mandatory" items for candidates with your Databricks depth. Apply now; don't wait to close every gap.

SECTION 11: Personal Brand Summary

Unique Value Proposition You're an Azure Databricks engineer who builds like an architect — delivering governed, scalable lakehouse platforms (medallion architecture, Delta Lake, Unity Catalog) in regulated, high-stakes environments, backed by six Azure and Databricks certifications and a proven SQL-Server-to-Databricks migration.

Professional Positioning Statement "I help organizations modernize their data platforms on Azure — turning fragmented, multi-source data into governed, reliable lakehouses that teams trust. I combine hands-on ETL/ELT engineering in Python and PySpark with architect-level design and a governance-first mindset shaped by years in regulated healthcare analytics."

Recruiter Headlines (pick per channel)

- **LinkedIn headline:** Azure Databricks Data Engineer | ETL/ELT • Delta Lake • Unity Catalog | Python • PySpark • SQL | 6x Azure/Databricks Certified
 - **Resume tagline:** Data Solution Architect & Azure Databricks Engineer — Scalable Lakehouse Pipelines, Governance, and Platform Modernization
 - **Elevator one-liner:** I build governed, scalable Azure data platforms — and I have the migrations and certifications to prove it.
-

SECTION 12: Interview Talking Points

Most Impressive Achievements

- **The MS SQL Server → Azure Databricks migration.** This is your headline story: end-to-end, medallion architecture, Delta Lake, Unity Catalog. Be ready to walk through the *why*, the *how*, and the *outcome*.
- **Metadata-driven ETL across four enterprise sources** (Epic, D365, Workday, Five9) — shows you can design reusable, scalable ingestion rather than one-off scripts.
- **Six Azure/Databricks certifications** — concrete, verifiable proof of breadth.
- **Career progression at EXL Service** from Data Analyst to Solution Architect.

Key Stories to Prepare (STAR format)

1. **Platform modernization:** The SQL-to-Databricks migration — Situation (legacy SQL Server limits), Task (modernize to lakehouse), Action (medallion + Delta + Unity Catalog, ADF pipelines), Result (governed, scalable platform).
2. **A data quality / governance challenge:** How you enforced HIPAA-aligned privacy and data quality — pick a concrete control you implemented.
3. **A multi-source integration problem:** Reconciling differing schemas/cadences across Epic, D365, Workday, Five9.
4. **A performance or cost-optimization win:** Tie in Azure FinOps and any SQL/Spark optimization you did (you optimized SQL data models in the Optum RCM work).
5. **Cross-functional / stakeholder story:** Translating business requirements into technical solutions (your offshore delivery leadership).

Areas Recruiters Are Likely to Probe (prepare answers in advance)

- **"Have you used Azure Synapse?"** — Be honest: position Databricks as your deep specialization and explain how the lakehouse model maps to Synapse concepts; mention you're actively learning it.
- **"Tell me about your Kafka / streaming experience."** — If you have batch-heavy experience, acknowledge it and discuss Structured Streaming + Kafka as a clear next step; show conceptual understanding.
- **"How do you implement cloud security?"** — Lean on your Azure Administrator cert: Azure AD, Key Vault secret scopes in Databricks, encryption, RBAC via Unity Catalog.
- **"This is a Data Engineer role and you're a Solution Architect — are you sure?"** — Reassure: you remain deeply hands-on (Python/PySpark/SQL daily), you value build-focused work, and the architect title reflects scope, not a move away from engineering.
- **Quantified impact** — your resume is light on hard numbers, so prepare 2–3 specifics *before interviews*: pipeline volumes, number of tables/sources migrated, runtime or cost improvements, dashboard adoption. Even approximate, defensible figures strengthen every story.

Prepared as a ready-to-use toolkit. Replace bracketed placeholders (names, dates) before sending. All content is drawn strictly from your resume — verify any figures you add are accurate before quoting them to employers.